

Search Algorithms

Song Liu (song.liu@bristol.ac.uk)

GA 18, Fry Building,

Microsoft Teams (search "song liu").

What is Search?


- Search algorithm finds object(s) from a **candidate set** according to a **searching criterion**.

Many CS problems are search problems:

- Find the index of the maximum (smallest) element in an array.
- Find the index of a specific element in a sorted array.
- Find x that maximizes a function $f(x)$.
- Find the best move in a chess game.
- Find the key to a password encryption.


Cryptography

- Cryptography is a classic search problem.
- To decrypt messages, you need to find "keys" from a huge candidate set: The mathematician (Alan Turing) who tried to solve this problem invented modern computers.



Machine Learning

- Many machine learning problems are search problems:
- Find a model from a **model family** that minimizes/maximizes objective functions:
 - A good weather prediction model should accurately predict temperatures with minimum error when comparing to the actual weather data.
 - A good advertisement model should recommend user ads that leads to maximum click rates.
- Commonly used model family (such as neural networks) usually a huge search space.



Search Problem

More formally,

$$\{i^*\} = \arg_{i \in \mathcal{F}} \min f(i)$$

where \mathcal{F} defines a search space and f defines the search criterion.

For example, finding the smallest element in a length n array:

$$\{i^*\} = \arg_{i \in \{1, \dots, n\}} \min s_i,$$

where s_i is the i -th element in the array s .

finding the index of the element equals 4 in a length n array:

$$\{i^*\} = \arg_{i \in \{1, \dots, n\}} \min |s_i - 4|,$$

Search Problem

When the search space is discrete, and is not large, we can enumerate the entire \mathcal{F} to find the best fit(s).

- Like what we have done in `find_min_idx` function.

However, what if the search space is infinite?

$$\{x^*\} = \arg_{x \in \mathbb{R}} \min f(x),$$

Do we search the entire real domain?

Greedy Algorithm

Greedy algorithm is the name of a **set of search algorithms**.

1. Greedy algorithm starts from a smaller, but more manageable **subproblem**.
2. Greedy algorithm finds the best possible solution (hence the name greedy) for the current subproblem.
3. From the solution, it **revises** the subproblem and solves a new subproblem.
4. Until a certain stopping criterion is satisfied.

For greedy algorithm to work, the subproblem must provide **enough information** to the original problem.

Hill Climbing:

Consider the problem of finding the minimizer of a function:

$$\{x^*\} = \arg_{x \in \mathbb{R}} \max f(x),$$

How do we split this problem into meaningful subproblems?

Hill Climbing:

We can start at an arbitrary location x_0 and solve for the following problem:

$$\{x_1^*\} = \arg_{x \in [x_0 - \epsilon, x_0 + \epsilon]} \max f(x),$$

where ϵ is a fixed value.

The next step, we solve

$$\{x_2^*\} = \arg_{x \in [x_1^* - \epsilon, x_1^* + \epsilon]} \max f(x)$$

and so on...

- After finding the maximizer of the previous subproblem, we restart search for the maximum centered around the previous maximizer.
- Hence the name "Hill Climbing".

Hill Climbing:

The algorithm stops when successive steps finds identical optimal solutions, i.e. $x_{t+1}^* = x_t^*$.


For each sub-problem,

$$\{x_{t+1}^*\} = \arg_{x \in [x_t^* - \epsilon, x_t^* + \epsilon]} \max f(x),$$

we can find an approximate solution via grid search:

1. Split $[x_t^* - \epsilon, x_t^* + \epsilon]$ into discrete "grid points".
 - e.g. $[-1, 1]$ can be discretized as $[-1, -.9 \dots .9, 1]$.
2. Search for the grid point that gives the smallest function value.

Hill Climbing:



The Subproblem

```
// This function returns xmax that maximizes f(x)
// between x-epsilon and x+epsilon
// complete the code yourself
double subproblem(double xt, double epsilon){
    double x = xt - epsilon;
    double fmax = -100;
    double xmax = x;
    while(x <= xt + epsilon){
        if(____){
            //fill out the if statement
        }
        x += .1;
    }
    return xmax;
}
```

Write the `main` function

```
void main(){
    double epsilon = .5;
    double x0 = 1;

    double xt = x0;
    double xt_1 = 100;

    while( fabs(xt - xt_1) >= 1e-5){
        xt_1 = xt;
        _____ // fill out the blank
        printf("f(%.4f) = %.4f\n", xt, f(xt));
    }

    printf("Maximum is at %f with value %f\n", xt, f(xt));
}
```

Output

$f(-3.5000) = -0.3430$

$f(-3.0000) = -0.2658$

$f(-2.5000) = -0.1380$

$f(-2.0000) = 0.0361$

$f(-1.5000) = 0.2489$

$f(-1.0000) = 0.4895$

$f(-0.5000) = 0.7448$

$f(-0.0050) = 0.9975$

$f(0.4900) = 1.2351$

$f(0.9850) = 1.4427$

$f(1.4850) = 1.6080$

$f(1.9850) = 1.7168$

$f(2.4850) = 1.7597$

$f(2.5400) = 1.7602$

$f(2.5400) = 1.7602$

Maximum is at 2.540000 with value 1.760173

Questions


Is this algorithm guaranteed to find the maximum of $f(x)$?

How will epsilon affect the performance of the algorithm?

Locally Optimum


- No, the algorithm is **not** guaranteed to find the maximum.
- Solving subproblems successively may not lead to the global optimal solution.
 - The subproblems may not provide necessary information about **the global structure** of $f(x)$ to reach the global optimal solution.
 - It may only reach a solution that is **locally optimal**, which means the solution is optimal only within **the reach of our search spaces**.

Local Optimum



Algorithm stuck at x_2^* !

Local Optimum



With a bigger ϵ , algorithm is **not** stuck at local optimal!

- However, bigger ϵ means solving the subproblem is more time-consuming (more grid points).
- We have to make the trade off between a better solution and computational time.

TicTacToe

- In some cases, the search space \mathcal{F} is finite and well-defined, but the searching criteria is not.
- Consider the game TicTacToe:

```
X O *  
* O *  
X X *
```

where `*` is an empty spot.

- Players `x` and `o` try to connect their pieces as a straight line, which can be a column, a row, or a diagonal line.

```
X O *  
* O O  
X X X // X wins!
```

TicTacToe

- Given a specific game, what is the optimal move?

```
X O *  
* X *  
* * *
```

- The "moves" are finite and well-defined (the empty spots on the board).
- However, how do we define "the optimal move"?
 - Sure, a good move leads to winning the game, but how to specify the optimality in an algorithmic manner?

TicTacToe

- We can use the greedy algorithm.
- Instead of targeting on "winning the game" in the long run, we focus on the **next step only**.
- Let the game board be a 3 by 3 matrix.

Recall the game is trying to connect pieces into a straight line.

- For a good move i, j ,
 - It should maximize the number of our own pieces on i -th row, j -th column and diagonal lines (if i, j is on the diagonal).
 - It should avoid opponent pieces along these lines, which would prevent us from forming a line.

Greedy Algorithm

At step t , the optimal move is determined by

$$\{(i_{t+1}^*, j_{t+1}^*)\} = \arg_{(i,j) \in \mathcal{F}_t} \min f(i, j),$$

- \mathcal{F}_t is the set of feasible moves.
- $f(i, j)$ scans i -th row, j -th column and diagonal lines (if i, j is on the diagonal or anti-diagonal).
- The function value is the number of self-pieces on these lines **minus** the number of opponent-pieces.

Greedy Algorithm

Suppose M is the game board. We can define

$$\begin{aligned} f(i, j) := & \sum_k \mathbb{1}(M_{i,k} = \text{self}) - \sum_k \mathbb{1}(M_{i,k} = \text{opp}) \\ & + \sum_k \mathbb{1}(M_{k,j} = \text{self}) - \sum_k \mathbb{1}(M_{k,j} = \text{opp}) \\ & + g(i, j) + h(i, j). \end{aligned}$$

$$g := \begin{cases} \sum_k \mathbb{1}(M_{k,k} = \text{self}) - \sum_k \mathbb{1}(M_{k,k} = \text{opp}), & i = j \\ 0, & i \neq j \end{cases}$$

$$h := \begin{cases} \sum_k \mathbb{1}(M_{k,2-k} = \text{self}) - \sum_k \mathbb{1}(M_{k,2-k} = \text{opp}), & i = 2 - j \\ 0, & i \neq 2 - j \end{cases}$$

The Game Class

```
class tictactoe{
    matrix board; // the game board
    int isplayable(int i, int j) { // is i,j playable?
        if (board.get_elem(i, j) == '*') {
            return 1;
        }
        return 0;
    }

public:
    tictactoe(): board(3,3){// how you initialize a field
        //TODO: initialize the board with *
    }
    void play(int i, int j, char player) {
        if(isplayable(i, j)) {
            // fill out the blanks
            board.set_elem(__, __, __);
        }
    }
};
```

- The tictactoe "has a" board.

The Helper Functions

Implement the following private functions

```
int pieces_at_row(int i, char player){
    //Count player's pieces at i-th row
    int count = 0;
    for(int j = 0; j < 3; j++){
        if(board.get_elem(i, j) == player) count++;
    }
    return count;
}

int pieces_at_col(int j, char player){
    //TODO: Count player's pieces at j-th column
}

int pieces_at_diag(char player){
    //TODO: Count player's pieces at diagonal line
}

int pieces_at_anti_diag(char player){
    //TODO: Count player's pieces at anti-diagonal line
}
```

Evaluate Moves

```
int evaluate(int i, int j, char player){
    char opponent = (player == 'X') ? 'O' : 'X';
    // evaluate move i,j
    // the higher the score is, the better the move is
    int score = 0;

    score += pieces_at_row(i, player);
    score -= pieces_at_row(i, opponent);
    score += pieces_at_col(j, player);
    score -= pieces_at_col(j, opponent);

    if(i == j){
        score += pieces_at_diag(player);
        score -= pieces_at_diag(opponent);
    }
    if(i + j == 2){
        score += pieces_at_anti_diag(player);
        score -= pieces_at_anti_diag(opponent);
    }

    return score;
}
```

Make the Move!

Write the following public function:

```
void play(char player){  
    // TODO: It play the next move for the "player",  
    // play the move (i,j) is the maximizer of f(i,j)  
  
}
```

Let it Play!

```
int main()
{
    tictactoe game;
    game.play(0, 0, 'X');
    game.play(1, 0, 'O');
    game.play(0, 2, 'X');
    game.play(1, 2, 'O');
    game.print();

    printf("AI plays... \n");
    game.play('X');
    game.print();
}
```

- You may need to implement your own `print` function to print out the game.

Output

```
X * X
```

```
O * O
```

```
* * *
```

```
AI plays...
```

```
X X X
```

```
O * O
```

```
* * *
```

We win!

However, is our strategy optimal?

A Suboptimal Move

```
X O *
```

```
X X O
```

```
* * *
```

AI plays...

```
X O O
```

```
X X O
```

```
* * *
```

- You can design a more complex subproblem anticipating your opponent's counter moves, but it would increase computational complexity too.
- We have to compromise between computational time and optimality of our solution.

A Min-Max Search

You can think about an optimal move:

- If there is a chance to win the game in the next round, take it.
- If there is no game-winning moves, take the move that minimizes our opponent chance to win the game.
- Our opponent will use **the same strategy** against us.

This is called **min-max algorithm**.

Conclusion

- Search Problem finds objects(s) from a search space given a searching criterion.
- Greedy algorithms
 - Hill Climbing
 - TicTacToe
- Both of them can find locally optimal solutions, but neither of them is guaranteed to find the global optimum.
- Greedy algorithm is myopic, which means it only focuses on optimizing the subproblem, which may NOT lead to **the global optimal solution.**